
GEN-AI
APPS
CREATOR
https://CloudAge.global

GEN-AI APPS CREATOR IN AWS

KASHIF JAMAL

Rapidly develop and deploy production-ready generative AI applications

APPLICATION COST BREAKDOWN

Service/Component	*Dimensions*	*Monthly Cost (USD)*
Lambda, S3, Parameter Store	*32 KB/message, 5 min/connection*	*Included*
CloudWatch	*1.5 GB logs (verbose experimentation)*	*$7.23*
DynamoDB	*Conversation/LLM tables, 2 GB total storage*	*$3.05*
Subtotal (Excluding LLMs)		***$10.89***
Bedrock (Nova Pro LLM)		
100 daily interactions	*90K input tokens/day = $0.152 × 30*	*16K output tokens/day = $0.0512 × 30*
Total (All Components)	*Use case + Bedrock (Nova Pro).*	***$17.00***

Notes:

- Cost breakdown assumes 100 user interactions per day.

- LLM costs are calculated separately and added to the core AWS service costs.

Overview

Building generative AI applications can be complex for teams that don't have deep AI expertise. Generative AI Application Builder on AWS simplifies this process, helping you develop, test, and deploy AI applications without extensive AI knowledge. This solution speeds up your AI development by easily incorporating your business data, comparing the performance of large language models (LLMs), running multi-step tasks through AI Agents, quickly building extensible applications, and deploying them with enterprise-grade architecture. Generative AI Application Builder comes with a ready-to-use generative AI chatbot and API that can be quickly integrated into your business processes or applications.

This solution includes integrations with Amazon Bedrock and its LLMs in addition to LLMs deployed on Amazon SageMaker. It uses Amazon Bedrock tools for Retrieval Augmented Generation (RAG) to enhance AI responses, Amazon Bedrock Guardrails to implement safeguards and reduce hallucinations, and Amazon Bedrock Agents to create workflows for complex tasks. You can also connect to other AI models using LangChain or AWS Lambda. Start with the simple, no-code wizard to build AI applications for conversational search, AI-generated chatbots, text generation, and text summarization.

Benefits

This solution allows users to experiment quickly by removing the heavy lifting required to deploy multiple instances with different configurations and compare outputs and performance. Experiment with multiple configurations of various LLMs, prompt engineering, enterprise knowledge bases, guardrails, AI agents, and other parameters.

With pre-built connectors to a variety of LLMs, such as models available through Amazon Bedrock, this solution gives you the flexibility to deploy the model of your choice, as well as the AWS and leading FM services you prefer. You can also enable Amazon Bedrock Agents to fulfill various tasks and workflows.

- **Built with AWS Well-Architected design principles, this solution offers enterprise-grade security and scalability with high availability and low latency, ensuring seamless integration into your applications with high performance standards.**

Extend this solution's functionality by integrating your existing projects or natively connecting additional AWS services. Because this is an open-source application, you can use the included LangChain orchestration layer or Lambda functions to connect with the services of your choice.

The Agent Use Case enables users to hand off tasks for completion using Amazon Bedrock Agents. You can select a model, write a few instructions in natural language, and Amazon Bedrock AI Agents will analyze, orchestrate, and complete the tasks by connecting to your data sources, or other APIs, to fulfill your request.

Architecture diagrams

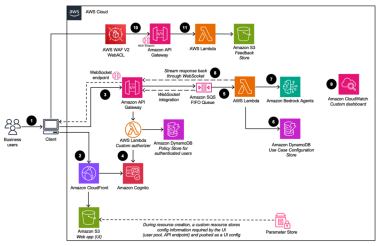

DEPLOY WITH CONFIDENCE

Everything you need to launch this AWS Solution in your account is right here
Ready to deploy?
Open the AMI in the AWS Console to begin setting up the infrastructure.

AWS Configure
bash deploy.sh